703650 VO Parallel Systems WS2020/2021
13 Dwarfs of HPC

Philipp Gschwandtner

Overview

- **13 Dwarfs of HPC**
  - abstract application categories

- **"Tales from the Proseminar"**
  - Daniel's Weird `int` Problem
  - g++ vs. gcc when compiling C code

Motivation

‣ MPI API or concepts such as "load imbalance" or "task-parallelism" are still pretty low-level characteristics of parallel programs

‣ we need to be able to recognize higher-level classes of applications and discuss them

‣ this lecture presents the most prominent classes of HPC applications

    ‣ many new applications you encounter will fit into these categories

How and Why are Dwarfs Defined?

▸ group applications by similarity in computation and data structures

  ▸ first published by Asanovic et al in *The Landscape of Parallel Computing Research: A View from Berkeley*

  ▸ https://www2.eecs.berkeley.edu/Pubs/TechRpts/2006/EECS-2006-183.pdf

▸ purely algorithmic, implementation-independent

  ▸ enables cross-platform reasoning and cross-application resource sharing (e.g. libraries)

▸ serve as small, abstract, high-level benchmarks for studying new

  ▸ programming models
  ▸ communication patterns
  ▸ hardware architectures, topologies
  ▸ …

▸ used to kick off innovation in all of these aspects

7 Original Dwarfs of HPC

▶ **What are we going do discuss?**

    ▶ 1. Dense Linear Algebra

    ▶ 2. Sparse Linear Algebra

    ▶ 3. Spectral Methods

    ▶ 4. N-body Methods

    ▶ 5. Structured Grids

    ▶ 6. Unstructured Grids

    ▶ 7. Monte Carlo Methods

▶ **What have you already heard?**

    ▶ matrix mul (first MPI lecture)

    ▶ particle-in-cell

    ▶ heat stencil

    ▶ Monte Carlo π

Dense Linear Algebra

▸ **data is stored in densely-populated matrices (or vectors)**

    ▸ data is stored uncompressed ("as is")

    ▸ data access via strides, often unit stride

▸ **e.g. matrix multiplication, LU decomposition, Gauss-Seidel, …**

▸ **rarely done manually, there are a TON of libraries out there**
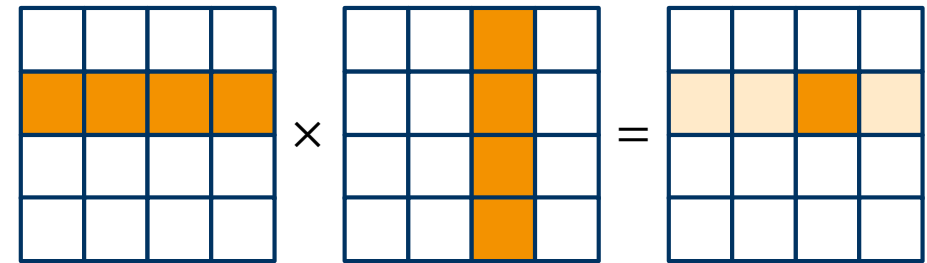
Dense Linear Algebra: Characteristics

- **naïve implementations often memory bound**
  - remember: memory wall!
  - caches and prefetching helps

- **simple but significant data structure**
  - stride often enables/prevents vectorization (SIMD)
  - column- vs. row-major access (often transposed) affects cache efficiency

- **still the default measure for performance in HPC**
  - e.g. TOP500 uses HPL, a high performance LINPACK benchmark
  - but nowadays not the only one (e.g. HPCG)

```c
for (int i = 0; i < N; ++i) {
  for (int j = 0; j < N; ++j) {
    double tmp = 0.0;
    for (int k = 0; k < N; ++k) {
      tmp += A[i][k] * B[k][j];
    }
    result[i][j] = tmp;
  }
}
```

```asm
vgatherqpd ymm0{k2}, [rax+ymm5*1]
vmulpd  ymm0, ymm0, YMMWORD PTR [rdx+rdi]
...
```

Dense Linear Algebra: Optimizations

▸ **loop blocking or tiling**
  ▸ do not work on single elements but smaller blocks (e.g. 2x2 or 32x32)
  ▸ exploits locality and cache
  ▸ also, lots of other HOTs (<u>H</u>igher <u>O</u>rder <u>T</u>ransformations)

▸ **vectorization (SIMD, e.g. SSE/AVX)**
  ▸ might entail modifications, e.g. transposing matrix A or B in matrix mul

▸ **hardware-specific instructions**
  ▸ e.g. fused multiply-add (FMA)

```cpp
for (int ii = 0; ii < N; ii += ib) {
 for (int jj = 0; jj < N; jj += jb) {
  for (int k = 0; k < N; ++k) {
   for (int i = ii; i < ii+ib; ++i) {
    for (int j = jj; j < jj+jb; ++j) {
     // ... process single tile
    }
   }
  }
 }
}
```

Sparse Linear Algebra

▶ **data is stored in sparsely-populated matrices (or vectors)**

  ▸ vast majority of data is zero

  ▸ data is stored in compressed format

  ▸ data often accessed indirectly via indices

▶ **e.g. conjugate gradient, Google's PageRank, data mining**
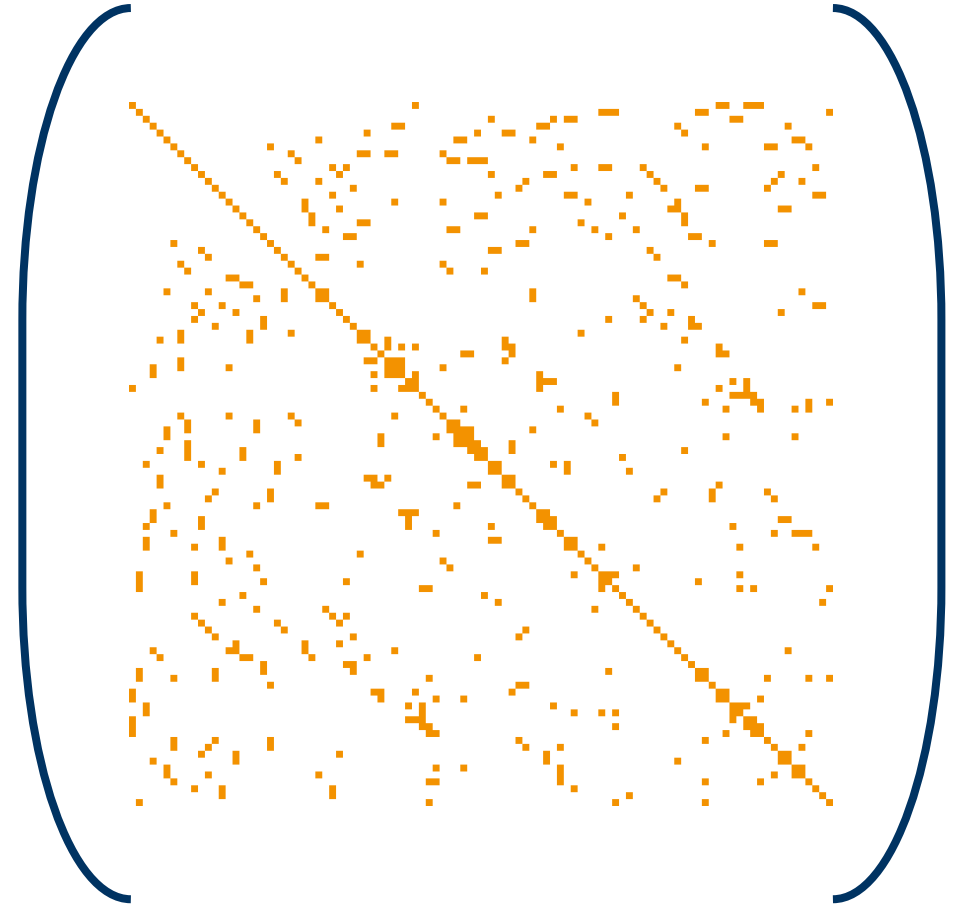
Sparse Linear Algebra: Characteristics

▸ **computationally or memory limited**

  ▸ depends on sparsity of data, data structure representation and algorithm

▸ **different data structures available**

  ▸ e.g. coordinate scheme (COO) or "triplet format" or similar: (i, j, a_{ij})

  ▸ array of structs vs. struct of arrays

  ▸ not necessarily sorted!

```c
typedef struct sparseElement {
    int i; int j; double value;
} sparseElement;
sparseElement sparseMatrix[SIZE];
sparseMatrix[0].i = 0;


// ############ vs. ############

typedef struct sparseMatrixT {
    int i[SIZE]; int j[SIZE];
    double values[SIZE];
} sparseMatrixT;
sparseMatrixT sparseMatrix;
sparseMatrix.i[0] = 0;
```

Sparse Linear Algebra: Optimizations

- **compressed row storage (CRS)**
  - two arrays of size N
  - one holds a_{ij}, the other j
  - third array points to start of row i in j
  - smaller memory footprint than COO
   - $2N + (m + 1)$ vs. $3N$

- **variants**
  - column-major variant (CCS)
  - store small blocks (2x2 or 4x4) instead of single elements, improves SIMDness
  - compressed diagonal storage (CDS)
   - remember: use domain-specific knowledge!

a_{ij}

2	-1	3	-1	2	1	0	3

j

0	1	0	1	2	1	2	3

start(i)

0	2	5	8

2	-1				
3	-1	2			
	1	0	3		
		4	-2	2	
			1	-3	1
				3	2

Spectral Methods

- **data in frequency domain, not time or space**
  - can include multiple stages of computations alternating between local and global communication

- **e.g. fast Fourier transform (FFT), audio and video signal processing**
  - e.g. "focus hunting" in contrast-based autofocus of photo/video cameras

Spectral Methods: Characteristics

▶ **usually implemented using butterfly patterns**

  ▶ multiple stages of multiply-add

  ▶ often latency limited due to global communication patterns (e.g. all-to-all)

▶ **often resemble structured or unstructured grid methods after transformation to frequency domain**
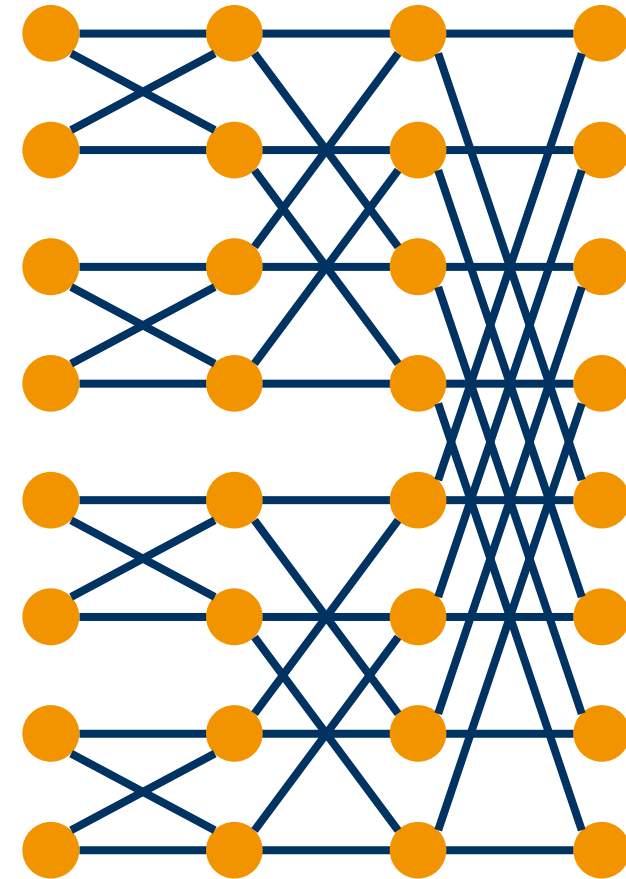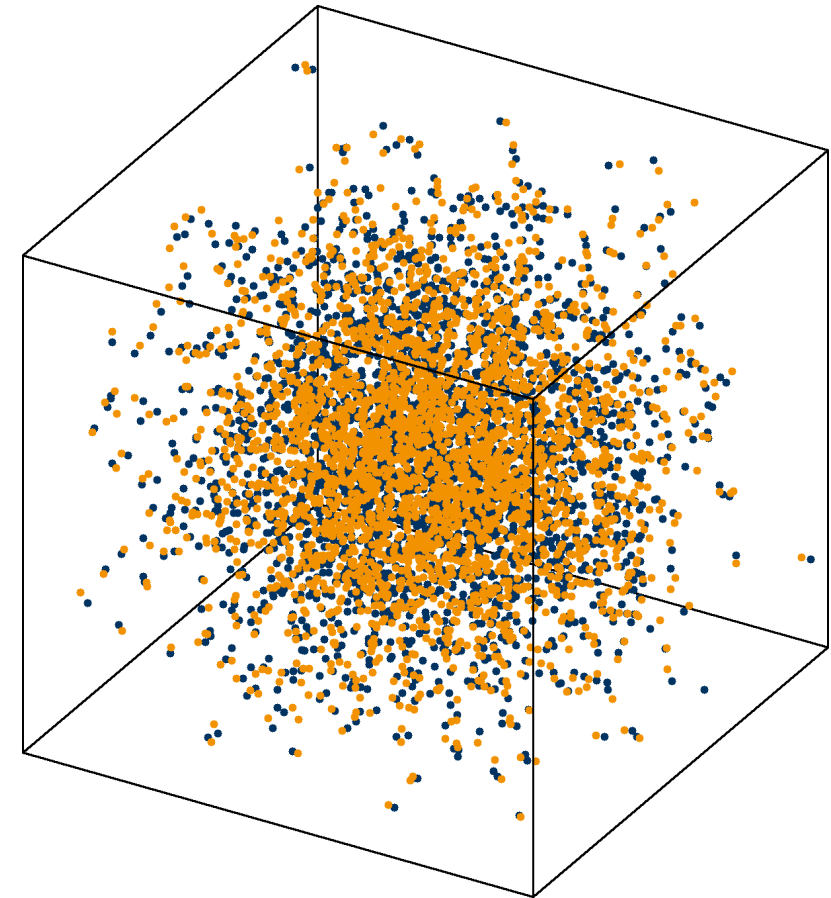
Spectral Methods: Optimizations

▸ **requires optimization of the transformation to frequency domain**
  ▸ relies on transposing data efficiently

▸ **afterwards, consider similar optimizations as for (un)structured grids**

▸ **severely restricted scalability to larger HPC systems**
  ▸ ongoing research
  ▸ lot's of it in math

N-body Methods

▶ **models interactions between discrete, moving points**

  ▶ often requires dynamic data structures

  ▶ varying spatial locality

  ▶ movement affects load balance and data access costs

▶ **e.g. galaxy collision simulations, molecular dynamics, protein folding**

N-body Methods: Characteristics

▸ don't want to spoil this week's assignment…

N-body Methods: Optimizations

▸ don't want to spoil this week's assignment…

Structured Grids

▸ **model interactions between discrete, fixed points**

　▸ grid structure described by pattern

　▸ topological information easily derived

　▸ usually high spatial locality

　▸ may be subdivided into finer grid ("adaptive")

▸ **e.g. heat transfer (stencil), computational fluid dynamics (CFD), octrees**

Structured Grids: Characteristics

▶ **structured nature is a key aspect**

  ▶ similar characteristics compared to dense linear algebra (e.g. row-major vs. column-major)

  ▶ memory addresses often predictable, facilitates e.g. prefetching

  ▶ adaptive grids possible

▶ **typically memory bound**

  ▶ e.g. 7-point stencil in 3D: load 7 data points for computing a single new one

  ▶ local communication only (ghost cell exchange with direct neighbor)

▶ **structured grid does not imply rectangular!**
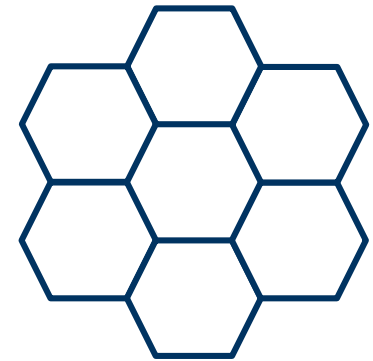
Structured Grids: Optimizations
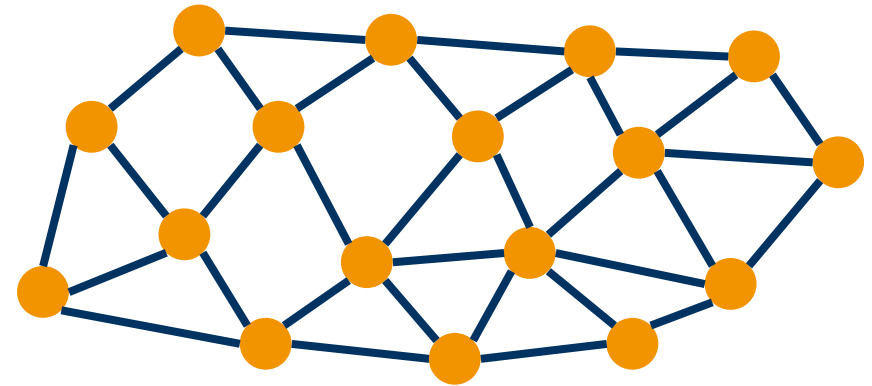
▸ **decomposition, decomposition, decomposition**

  ▸ remember in 3D: slabs, poles, cubes

  ▸ highly dependent on type of grid and specific use case

▸ **structured nature of the problem makes analytic prediction possible**

  ▸ performance models and simulators for simple stencil kernels (e.g. Kerncraft)

L1 misses: #
L2 misses: #
L3 misses: #
        …

Unstructured Grids

▶ **model interactions between discrete, fixed points**

- ▶ grid pattern described explicitly by individual connections
- ▶ irregular geometry and topology
- ▶ usually involves multiple levels of indirection when accessing data

▶ **e.g. computational fluid dynamics (CFD)**

Unstructured Grids: Characteristics

▶ **usually heavily latency bound due to indirect access**

   ▶ `… = cells[neighbors[i]]`

   ▶ `… = cell.getNeighbor(i)`

   ▶ also known as "pointer chasing"

▶ **similar problems compared to structured grids, e.g.**

   ▶ domain decomposition / adaptivity

   ▶ topological information

   ▶ ghost cell exchange

   ▶ but hardly analytically predictable

NUMECA
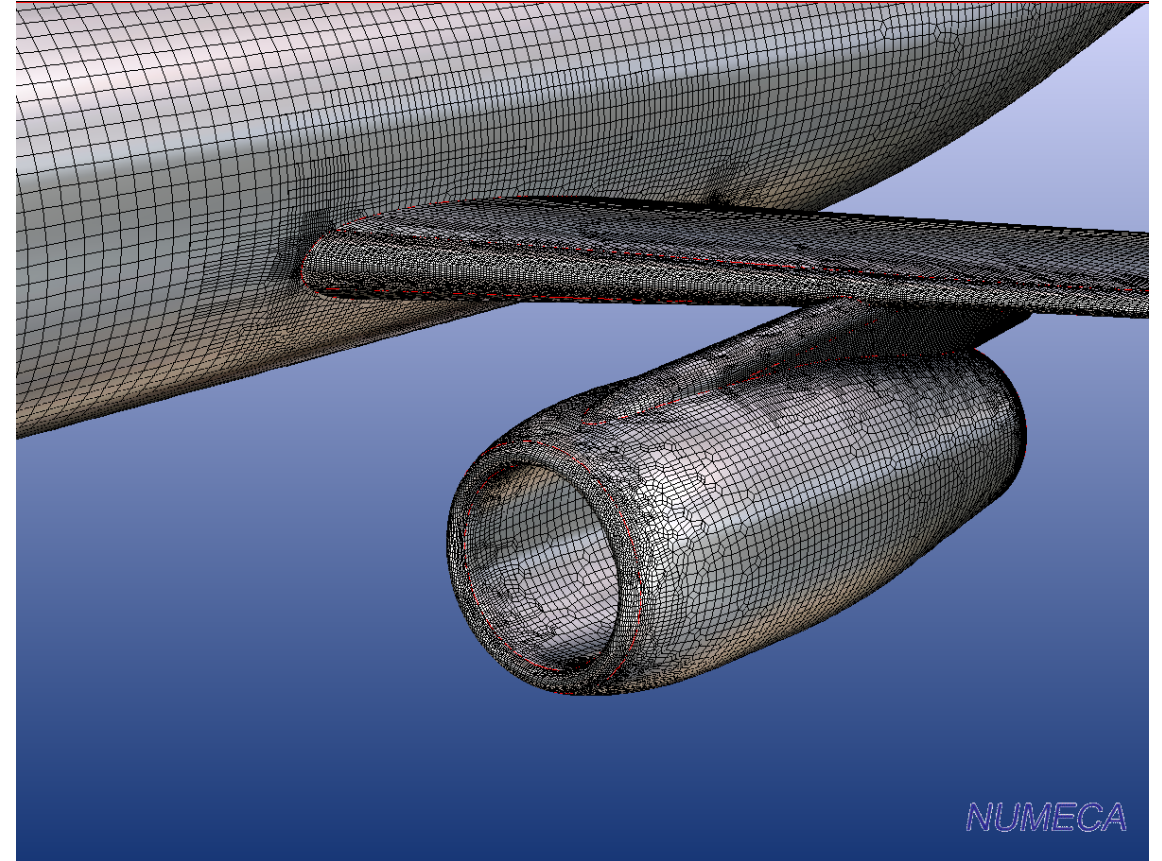
Unstructured Grids: Optimizations

▸ **problem space discretization and topology**

  ▸ which types of grid elements?

  ▸ which types of connections?

  ▸ cells, faces, vertices, edges, …

  ▸ should be efficient to load/store, navigate, and compute

▸ **decomposition, decomposition, decomposition**

  ▸ efficient ghost cell exchange required

  ▸ efficient grid navigation e.g. to access neighbors
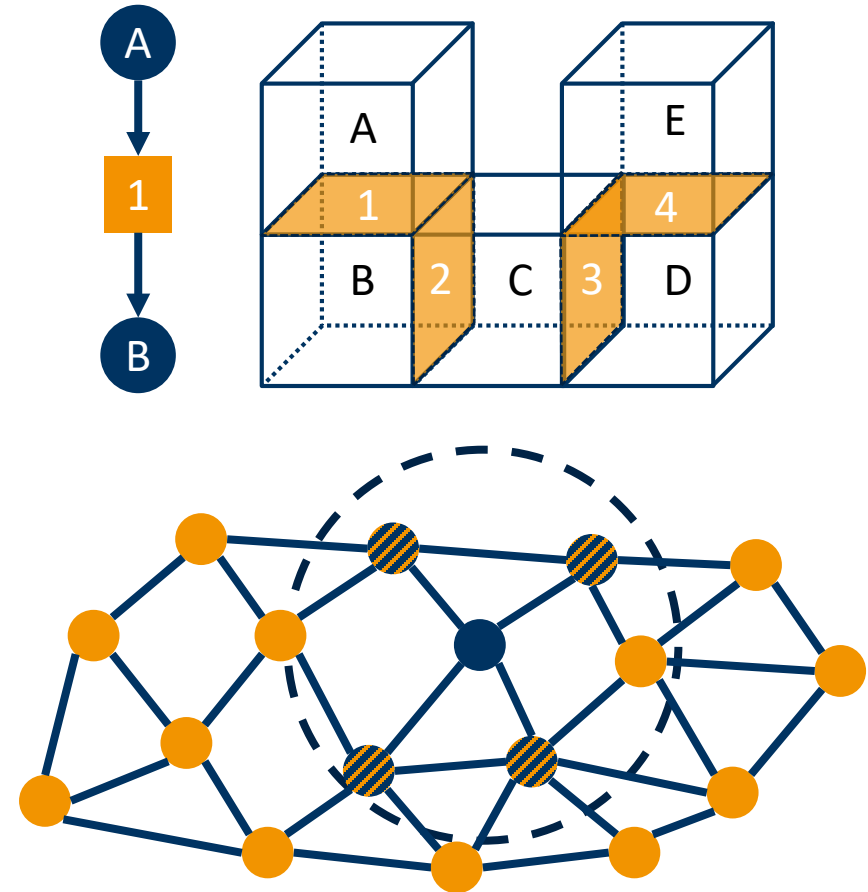
Monte Carlo Methods

▶ **also known as map-reduce**
  ▶ process data independently and merge the results

▶ **models statistical evaluation of repeated random trials**
  ▶ communication usually insignificant
  ▶ embarrassingly parallel, multiple copies of sequential method

▶ **e.g. numerical integration, quantum many-body problems, ray tracing**
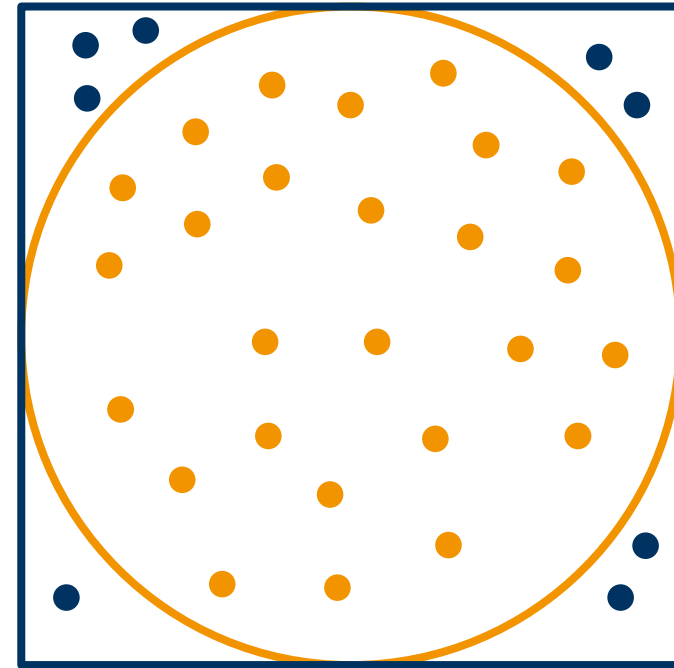
Monte Carlo Methods: Characteristics

▶ **parallelization almost a no-factor**

  ▶ similar to multiple sequential programs sharing some resources (e.g. L3 cache, random number generator)

  ▶ relatively inexpensive reduction

▶ **depends heavily on sequential optimizations**

Monte Carlo Methods: Optimizations

- **not much to do beyond sequential optimization**
  - ILP
  - prefetching
  - vectorization
  - resource contention (read: fast random number generation)

- **consider different hardware**
  - GPUs
  - FPGAs
  - ...

- **try to decrease cost of evaluating a sample**
  - increase number of samples if required

Additional Dwarfs

Additional Dwarfs

- ▸ 8. Combinational Logic
- ▸ 9. Graph Traversal
- ▸ 10. Dynamic Programming
- ▸ 11. Backtrack & Branch+Bound
- ▸ 12. Graphical Models
- ▸ 13. Finite State Machine

- ▸ slightly different focus compared to first 7 dwarfs
  - ▸ more (but not exclusively) on integer-heavy applications and machine learning
  - ▸ less on physical processes

Combinational Logic

▸ **generally involves performing simple operations on large amounts of integer data**
   ▸ e.g. computing cyclic redundancy codes (CRC)

▸ **often parallelizable on multiple levels**
   ▸ bit-level parallelism
   ▸ block-level parallelism

```c
uint8_t compute(uint8_t const msg[], int n) {
    uint8_t rem = 0;
    for (int byte = 0; byte < n; ++byte) {
        rem ^= (msg[byte] << (WIDTH - 8));
        for (uint8_t bit = 8; bit > 0; --bit) {
            if (rem & TOPBIT) {
                rem = (rem << 1) ^ POLYNOMIAL;
            } else {
                rem = (rem << 1);
            }
        }
    }
    return (rem);
}
```

Graph Traversal

▸ **traverse a number of objects in a graph and examine characteristics**

  ▸ e.g. searching, sorting, collision detection, decision trees, …

  ▸ usually heavy on data reads and lookups, very little computation and output

▸ **parallelizable over different paths in the graph**

  ▸ but indirect accesses are heavily latency-bound (c.f. unstructured grids)
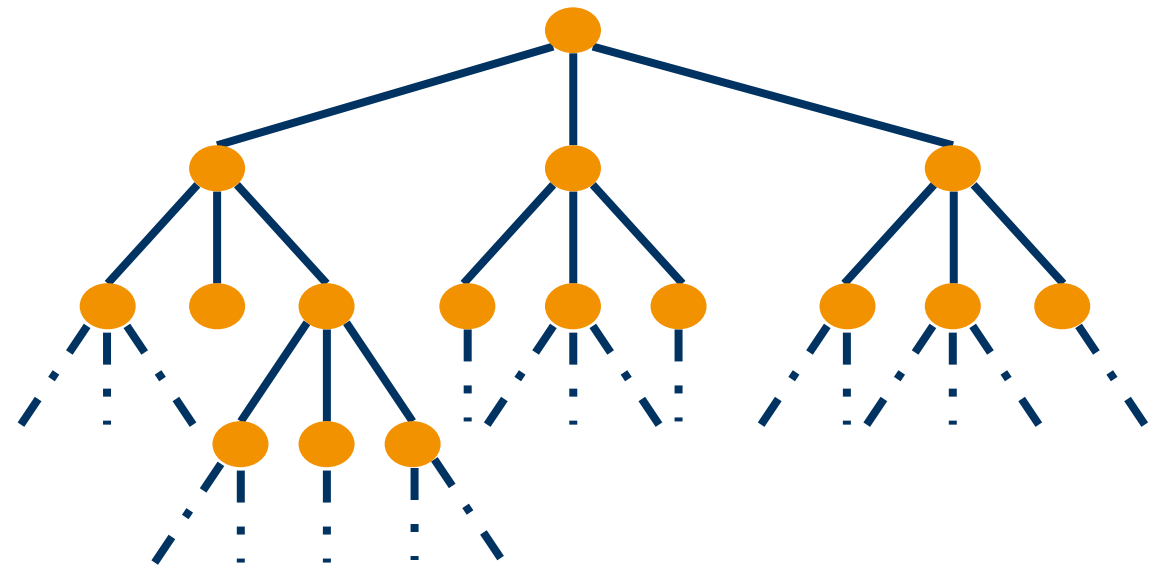
Dynamic Programming

▸ **method of computing solutions by solving simpler, overlapping sub-problems**

   ▸ applicable to problems where optimal result is composed of optimal results of sub-problems

   ▸ e.g. matrix-chain-multiplication

▸ **usually based on memoization**

   ▸ solve each sub-problem exactly once

   ▸ store and re-use the result

	A	B	C	D	X
	5x5	5x4	4x8	8x2	5x2
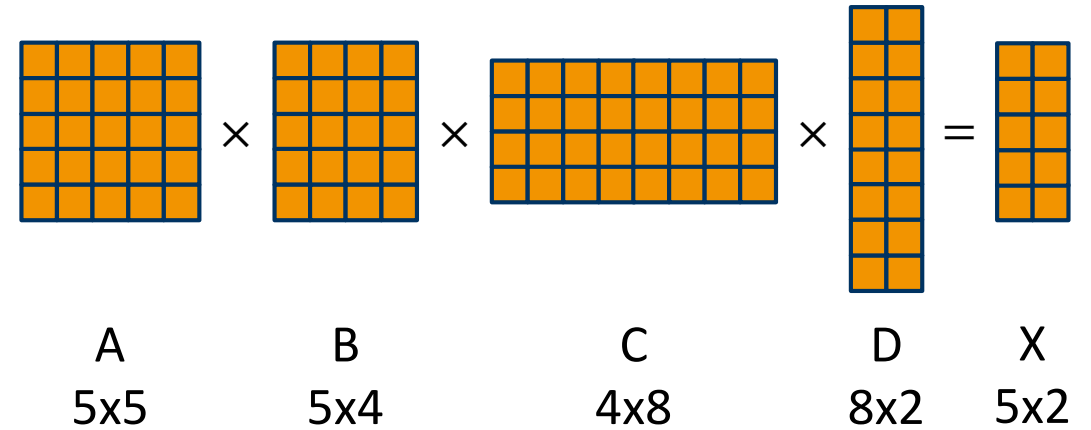
$$A \times \big(B \times (C \times D)\big) = X \quad 154 \text{ ops}$$

$$(A \times B) \times (C \times D) = X \quad 204 \text{ ops}$$

$$\big((A \times B) \times C\big) \times D = X \quad 340 \text{ ops}$$

Backtrack & Branch+Bound

▶ **search and optimization problems for very large problem spaces**

  ▸ incrementally build solution but discard if determined unsuitable

  ▸ e.g. n-queens problem

▶ **use divide & conquer strategy:**

  ▸ break down complex problem into smaller sub-problems until they become solvable

  ▸ solve sub-problems in parallel

Graphical Models

- also known as probabilistic models

- represent graphs consisting of random variables as nodes and dependencies as edges
  - e.g. Bayesian networks, Hidden Markov models

- ongoing research in math and computer science regarding parallelization and optimization

sun	rain
0.40	0.60

good	bad
0.90	0.10

weather

COVID19 news

mood

	good	bad
sun, good	0.95	0.05
sun, bad	0.70	0.30
rain, good	0.75	0.25
rain, bad	0.10	0.90
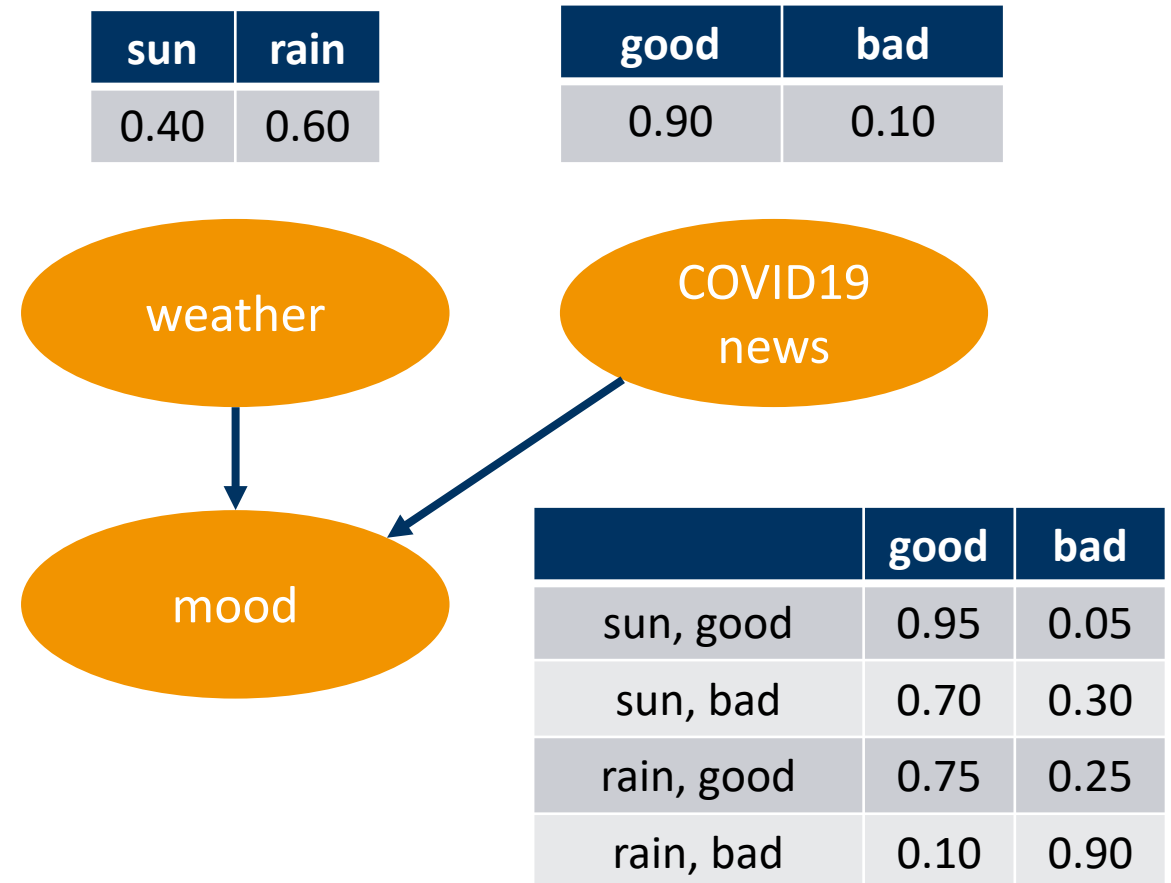
Finite State Machines

- represent interconnected set of states to be moved among
  - e.g. parsers

- can sometimes be decomposed into multiple state machines that act in parallel
  - ongoing research

S_0

(start) (start)

S_1 S_2

(start)

(create) (create)

S_3 S_4

(init) (sync) (init) (spawn)

S_5 (replicate) S_6 S_7

(sync) (start)

(replicate) (migrate)

S_8 S_9

(spawn) (init) (sync)

(sync)

… … …
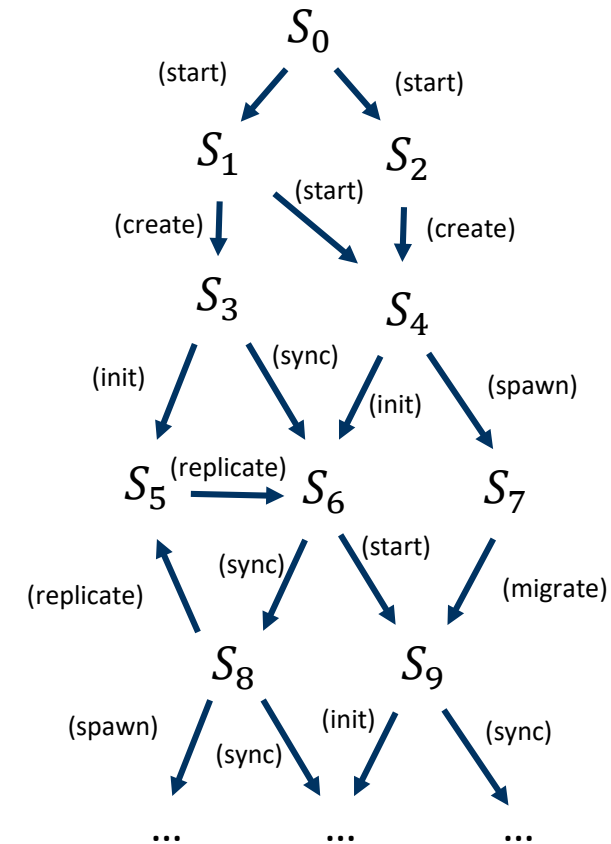
Literature Material

▶ White Paper by Berkeley University: „The Landscape of Parallel Computing Research: A View from Berkeley" from 2006: https://www2.eecs.berkeley.edu/Pubs/TechRpts/2006/EECS-2006-183.pdf

▶ Holds more detailed descriptions and related aspects

Note the Focus on Scientific Problems

▸ **MPI can also be used to implement**

  ▸ distributed-memory runtime systems

  ▸ emulate shared memory runtime systems on distributed memory (e.g. PGAS)

  ▸ provide the connecting parallelism layer for shared-memory or sequential systems

    ▸ e.g. use multiple accelerators in separate compute nodes (Celerity project @ UIBK)

    ▸ extend shared memory parallelism to distributed memory (MPI+X)

    ▸ …

▸ **still, the majority of codes is of scientific computing nature**

Tales from the Proseminar

Tales from the Proseminar: Daniel's Weird `int` Problem

▶ **sequential 2D heat stencil**

  ▸ 100x100 problem size

  ▸ -O0 (also with –O2) on LCC2

  ▸ gcc 4.8.5 (but also 9.2 and MSVC 2019)

  ▸ switching from `long long` loop iterators to `int` **halves** execution time

```
$ /usr/bin/time –f %E ./stencil_long 100 100
0:04.60
$ /usr/bin/time –f %E ./stencil_int 100 100
0:02.31
```

▶ **What the heck is going on?**

Tales from the Proseminar: Daniel's Weird `int` Problem

▸ profiled with `gprof` to find "hot spots", left is `int`, right is `long long`

```
%     cumulative    self                    %     cumulative    self
time   seconds    seconds       name        time   seconds    seconds       name
32.92    0.76      0.76       int.c:93      36.84    1.69      1.69     long.c:79
20.06    1.22      0.46       int.c:78      31.61    3.15      1.45     long.c:78
14.83    1.56      0.34       int.c:79      16.35    3.90      0.75     long.c:93
10.47    1.80      0.24       int.c:87       5.34    4.15      0.25     long.c:86
 9.59    2.02      0.22       int.c:86       3.27    4.30      0.15     long.c:87
```

```
74    // get temperature at current position
75    value_t tc = A[i];
76
77    // get temperatures of adjacent cells
78    value_t txl = (i % Nx != 0) ? A[i - 1] : tc;
79    value_t txr = (i % Nx != Nx - 1) ? A[i + 1] : tc;
      // ...... snip ......
90    if (Ny > 1)
91      B[i] = tc + 0.165 * (txl + txr + tyl + tyr + tzl + tzr + (-6 * tc));
92    else
93      B[i] = tc + 0.2 * (txl + txr + tzl + tzr + (-4 * tc));
94    // if ((int)B[i] < (int)A[i])
```

Tales from the Proseminar: Daniel's Weird `int` Problem

▶ far-fetched idea, but maybe branch (miss-)predictions? compared with `perf stat`:

```
  2328.650634 task-clock:u (msec)  #    0.996 CPUs
            0 context-switches:u   #    0.000 K/sec
            0 cpu-migrations:u     #    0.000 K/sec
          186 page-faults:u        #    0.080 K/sec
5,739,935,525 cycles:u             #    2.465 GHz
10,671,532,880 instructions:u      #    1.86  IPC
1,196,623,336 branches:u           # 513.870 M/sec
    1,030,678 branch-misses:u      #    0.09%


  2.338387625 seconds time elapsed
```

```
  4639.728721 task-clock:u (msec)  #    0.998 CPUs
            0 context-switches:u   #    0.000 K/sec
            0 cpu-migrations:u     #    0.000 K/sec
          186 page-faults:u        #    0.040 K/sec
11,444,184,736 cycles:u            #    2.467 GHz
10,972,976,005 instructions:u      #    0.96  IPC
1,196,977,569 branches:u           # 257.984 M/sec
    1,030,347 branch-misses:u      #    0.09%


  4.650327844 seconds time elapsed
```

Tales from the Proseminar: Daniel's Weird `int` Problem cont'd

```
value_t txl = ( i % Nx != 0 ) ? A[i-1] : tc;
```

```
mov     eax, DWORD PTR [rbp-20]
cdq
idiv    DWORD PTR [rbp-24]
mov     eax, edx
test    eax, eax
je      .L2
mov     eax, DWORD PTR [rbp-20]
cdqe
```

```
sal     rax, 3
lea     rdx, [rax 8]
mov     rax, QWORD PTR [rbp-32]
add     rax, rdx
movsd   xmm0, QWORD PTR [rax]
jmp     .L3
.L2: ...
.L3: ...
```

Tales from the Proseminar: Daniel's Weird `int` Problem cont'd

▶ Found instruction information on Agner Fog's blog for Intel Skylake architecture:

  ▶ https://www.agner.org/optimize/instruction_tables.pdf

▶ if you want to study compiler output: https://godbolt.org/

inst.	operands	μops	latency
idiv	r32	10	26
idiv	r64	57	42-95

Tales from the Proseminar: Daniel's Weird `int` Problem cont'd

▸ **root cause of the issues: single loop for multi-dimensional problem space**

   ▸ requires % operator to get boundaries (which are strided in linearized space)

   ▸ replace with loop nest and comparison operators (>, <, ==, !=)

▸ **potentially premature optimization and violates step 1 of** *"Four Steps to Creating an Optimized Parallel Program"*

Tales from the Proseminar: g++ vs. gcc for C Code

```c
void foo() {
    for(int i = 0; i < 10; ++i) { }; return;
}
```

g++   gcc

```asm
        push    rbp
        mov     rbp, rsp
        mov     DWORD PTR [rbp-4], 0
.L3:
        cmp     DWORD PTR [rbp-4], 9
        jg      .L2
        add     DWORD PTR [rbp-4], 1
        jmp     .L3
.L2:
        nop
        pop     rbp
        ret
```

```asm
        push    rbp
        mov     rbp, rsp
        mov     DWORD PTR [rbp-4], 0
        jmp     .L2
.L3:
        add     DWORD PTR [rbp-4], 1
.L2:
        cmp     DWORD PTR [rbp-4], 9
        jle     .L3
        nop
        pop     rbp
        ret
```

Summary

▸ **13 Dwarfs of HPC**

  ▸ abstract application categories

  ▸ facilitate cross-platform reasoning and cross-application component reuse

  ▸ 7 older ones, most well-studied physics problems

  ▸ 6 newer ones, partially of theoretical nature, subject to ongoing research

  ▸ give a broad perspective on HPC potential and limitations

▸ **"Tales from the Proseminar"**

  ▸ Daniel's Weird `int` Problem: Don't consider all `int` instructions to be cheap!

  ▸ g++ vs. gcc when compiling C code

Image Sources

- Dwarfs: http://pngimg.com/download/47261, http://rpgmke.wikidot.com/moonshae-isles-campaign

- CRC: https://barrgroup.com/Embedded-Systems/How-To/CRC-Calculation-C-Code

- Unstructured Mesh: https://resourcearea.cpu-24-7.com/en/numeca_welcome